

```

#define DIRECT_BLOCKS_COUNT 123
#define INDIRECT_BLOCKS_PER_SECTOR 128

/* On-disk inode.
| | Must be exactly BLOCK_SECTOR_SIZE bytes long. */
// 1 inode_disk == 1 block sector size i.e. 512 bytes long
// inode_disk == direct blocks + indirect blocks (disk addrs)
// file inode_disk == direct blocks + indirect blocks (disk addrs) of 512 bytes file chunk
// dir inode_disk == direct blocks + indirect blocks (disk addrs) of 64 dir entries (in 512 bytes)
// dir entry == 8 bytes storing disk addr of either subdir / file inode_disk
struct inode_disk
{
    // Data sectors -> 500 bytes
    // must be pointing at 1 block sector, could be:
    // 1. dir inode disk 2. file inode disk 3. indirect inode disk
    // 3. block size of dir entires 4. block size file chunk
    block_sector_t direct_blocks[DIRECT_BLOCKS_COUNT]; // 4*123 == 492 bytes
    block_sector_t indirect_block; // 4 bytes
    block_sector_t doubly_indirect_block; // 4 bytes

    // remaining part -> 12 bytes
    bool is_dir;
    off_t length; // File size in bytes. */
    unsigned magic; // Magic number. */
}; // total == 512 bytes

/* In-memory inode. */
struct inode
{
    struct list_elem elem; // Element in inode list. */ // for loop able
    int open_cnt; // Number of openers. */
    bool removed; // True if deleted, false otherwise. */
    int deny_write_cnt; // 0: writes ok, >0: deny writes. */
    block_sector_t sector; // inode_disk's disk addr / sector no.
    struct inode_disk data; // 1 block size 512 bytes of indexes to disk addrs
};

```

```

// file in RAM == directly pointing at 123 * 512 bytes of file at disk
// dir in RAM == directly pointing at 7872 dir entries at disk
// inode == dir inode_disk (supposedly in disk) in RAM
// dir inode_disk == direct blocks + indirect blocks (disk addrs) of block of 64 dir entries
// direct_block == points at 64 dir entries on 1 block sector (512 bytes)
// direct block alone == could support 123 * 64 = 7872 file / subdir's inode_disk
struct dir // memory
{
    struct inode *inode;           /* Backing store. */
    off_t pos;                     /* Current position. */
};

// dir version of file chunk, BUT divided in 64 small pieces n storing subdir / file inode_disk
// dir entry == 1 entry in a directory e.g. file or subdir
struct dir_entry // disk
{
    block_sector_t inode_disk_diskaddr; // pointing at either subdir / file inode_disk
    char name[NAME_MAX + 1];           // sub_dir name or file name
    bool in_use;                        /* In use or free? */
};

```

FS

PA (kernel stack)

disc

thread → fd-list[]

list of
fd[]

file_desc① → file
→ id → dir
fd②
fd③

root-dir - inode
dir - pos

root-dir - disk
addr of
inode
- inode-disk inode-disk

sub-dir - inode
dir - pos

sub-dir - disk
inode of
inode-disk

file - inode
- pos

file - disk addr
of
inode
- inode-disk

x 512
bytes

Bitmap

root-dir

inode-disk

root-dir

dir-entries.

Sub-dir

inode-disk

sub-dir

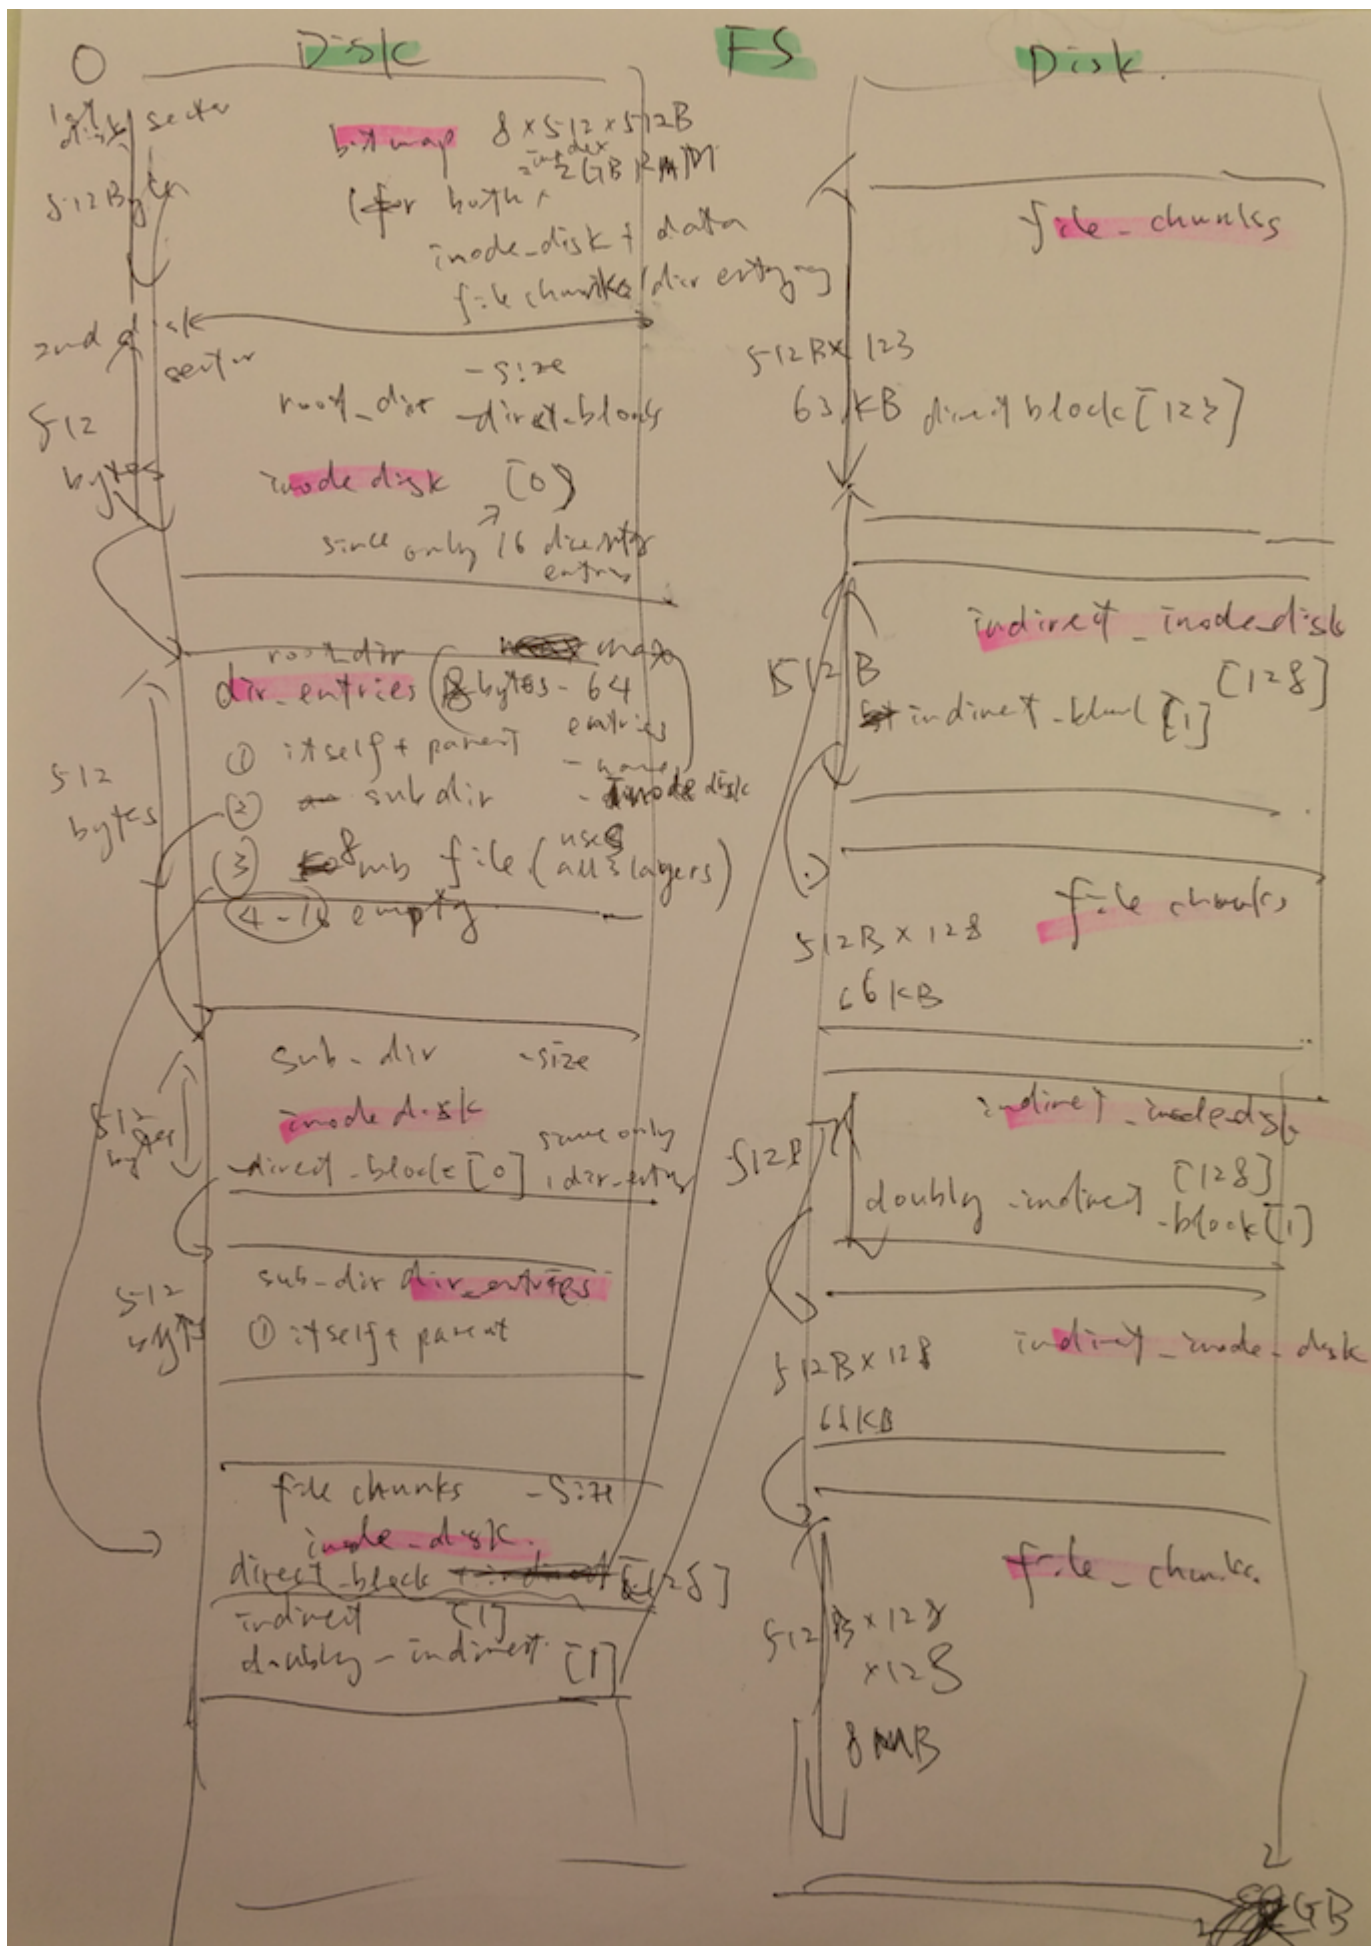
dir-entries.

file-chunks

inode-disk

file-chunks

✓ 512 bytes
except
file-chunk



DEMO inode on disk

- file -> file_inode (RAM) -> file_inode_disk (both RAM n disk) -> file_chunks (disk)

- `dir` -> `dir_inode` (RAM) -> `dir_inode_disk` (both RAM n disk) -> `dir_entries` (disk) -> `file_inode_disk/dir_inode_disk` (both RAM n disk)

HOW DATA ON DISK CHANGES BY KERNEL FUNCTIONS

Initial Disk State

`init_root_dir()`

`inode_create(sector, entry_cnt * sizeof(dir_entry))` -> POPULATE, WRITE `root_dir inode_disk` INTO DISK

1. `w()` `root_dir inode_disk` into inodes block 0
2. `r()` `w()` bitmap for `root_dir's dir_entries`
3. `w()` inodes block 0 to point to `data_block 0`
4. `w()` 0s at `data block 0 (root_dir dir_entries)`

`struct dir *dir = dir_open(inode_open(sector))` -> WRITE PARENT DIR + SUB DIR `dir_entries`

1. `for_loop()` `open_list` check `inode_disk` cache
2. if NO cache, `r()` `inode_disk` from `inode_disk` block 1
3. `populate()` `dir w inode_disk`

`struct dir_entry de.inode_sector = sector;` -> POPULATE "." ".." `dir_entry`

`inode_write_at(dir->inode, &de, sizeof de, 0)` -> WRITE `dir_entry` TO DISK

`dir_entries` at `data block 0`, w 2 `dir_entry` inside

`root_dir` at `data block 0` w 2 `dir_entries` : {`curr_dir`, `inode block 0`},{`parent_dir`, 0}

`inode bitmap` 10000000

`inodes` [d a:0 r:2] [] [] [] [] [] [] []

`data bitmap` 10000000

`data` [(.,0) (.,0)] [] [] [] [] [] [] []

DEMO `init_root_dir` (`ROOT_DIR_SECTOR`, 16)

- this is showing layout of hard disk
- bitmap
 - easy way of showing whether 1 block / disk / sector (512 bytes) is free
 - bitmap itself is 1 block / disk / sector (512 bytes) long
 - inode, data bitmap are combined into one single bitmap in pintos
- []
 - 1 block / disk / sector (512 bytes) long on disk
- inodes
 - file / dir `inode_disk`
 - d: dir `inode disk`, f : file `inode disk`, a: disk addr, r : number of files inside
 - 1 block / disk / sector (512 bytes) long
- data
 - 1 file chunk / 64 `dir_entries`
 - 1 block / disk / sector (512 bytes) long

```

creat("/y");

sys_create("/y") file

struct dir *parent_dir = dir_open_path(dir_path)    -> SEARCH FOR PARENT DIR INODE
1. read() root_dir inode_disk
2. traverse() dir_entries data block for matching dir name
3. if found dir_entry{name, inode}, populate() parent_dir w inode

free_map_allocate(1, &inode_sector)    -> BITMAP ALLOCATE file's inode_disk
1. r() w() bitmap for file's inode_disk

inode_create(inode_sector, initial_size, is_dir) -> POPULATE, WRITE file inode_disk INTO DISK
1. w() file inode_disk into inodes block 1
2. r() w() bitmap for file's file_chunks data block
3. w() inodes block 1 to point to data_block 1
4. w() 0s at data block 1 (reserve file_chunks data)

dir_add(dir, dir_name, inode_sector, is_dir) -> WRITE PARENT DIR dir_entries
1. w() file inode_disk addr at parent_dir dir_entries at data block 0

BEFORE:
inode bitmap  10000000
inodes        [d a:0 r:2] [] [] [] [] [] [] []
data bitmap   10000000
data          [(.,0) (.,0)] [] [] [] [] [] []

AFTER:
inode bitmap  11000000    -> free_map_allocate(size, &disk_addr_allocated) to store inode_disk
inodes        [d a:0 r:3] [f a:1 r:1] [] [] [] [] []    -> PintOS does not update dir size +
inode_create()
data bitmap   11000000    -> inode_reserve()
data          [(.,0) (.,0) (y,1)] [00000000] [] [] [] [] []    -> dir_add() + inode_reserve()

```

DEMO create file

- inodes == inode disk
 - 1st dir inode disk
 - 0 data disk addr
 - dir entries w 3 files / sub dirs
 - 1st file inode disk
 - 1 data disk addr
 - this file only takes up 1 file chunk
- data == dir entries / file chunks
 - 1st dir entry
 - 1 parent dir at 0 inode disk addr
 - 1 cur dir at 0 inode disk addr
 - 1 file "y" at 1 inode disk addr (you need to read from here to get to 1st file inode disk)
 - 1st file chunk
 - 1 file chunk filled with 0s

```

mkdir("/f");      -> filesystem_create() sub_dir "f"

struct dir *parent_dir = dir_open_path(dir_path)      -> SEARCH FOR PARENT DIR INODE
1. read() root_dir inode_disk
2. traverse() dir_entries data block for matching dir name
3. if found dir_entry{name, inode}, populate() parent_dir w inode

free_map_allocate(1, &inode_sector)      -> BITMAP ALLOCATE dir's inode_disk
1. r() w() bitmap for sub_dir's inode_disk

inode_create(inode_sector, initial_size, is_dir) -> POPULATE, WRITE sub_dir inode_disk INTO DISK
1. w() sub_dir inode_disk into inodes block 3
2. r() w() bitmap for sub_dir's dir_entries
3. w() inodes block 3 to point to data_block 2
4. w() 0s at data block 2 (sub_dir f's dir_entries)

dir_add(dir, dir_name, inode_sector, is_dir) -> WRITE PARENT DIR + SUB DIR dir_entries
1. w() sub_dir inode_disk addr at parent_dir dir_entries at data block 0
2. w() "." ".." inode_disk addr at sub_dir dir_entry at data block 2

BEFORE:
inode bitmap  11100000
inodes        [d a:0 r:5] [f a:1 r:1] [f a:-1 r:1] [] [] [] [] []
data bitmap   11000000
data          [(.,0) (.,0) (y,1) (z,2)] [u] [] [] [] [] [] []

AFTER:
inode bitmap  11110000      -> free_map_allocate(size, &disk_addr_allocated) to store inode_disk
inodes        [d a:0 r:6] [f a:1 r:1] [f a:-1 r:1] [d a:2 r:2] []
data bitmap   11100000      -> inode_reserve()
data          [(.,0) (.,0) (y,1) (z,2) (f,3)] [u] [(.,3) (.,0) 000000] [] ...

```

DEMO create dir

- inodes == inode disk
 - 1st dir inode disk
 - 0 data disk addr
 - dir entries w 6 files / sub dirs
 - 1st file inode disk
 - 1 data disk addr
 - this file only takes up 1 file chunk
 - 2nd file inode disk
 - -1 data disk addr (file was deleted ??)
 - this file only takes up 1 file chunk
 - 2nd dir inode disk
 - 2 data disk addr
 - dir entries w 2 files / sub dirs
- data == dir entries / file chunks
 - 1st dir entry
 - 1 cur dir at 0 inode disk addr
 - 1 par dir at 0 inode disk addr
 - 1 file "y" at 1 inode disk addr (you need to read from here to get to 1st file inode disk)
 - 1 file "z" at 2 inode disk addr (you need to read from here to get to 2nd file inode disk)
 - 1 dir "f" at 3 inode disk addr (you need to read from here to get to 2nd dir inode disk)
 - 1st file chunk
 - 1 file chunk with "u" in it
 - 2nd dir entry
 - 1 cur dir at 3 inode disk addr
 - 1 parent dir at 0 inode disk addr

```

fd=open("/y", O_WRONLY|O_APPEND); write(fd, buf, BLOCKSIZE); close(fd);

int fd = sys_open("/y")          -> open both file n parent_dir file is in

struct dir *dir = dir_open_path(dir_path)    -> SEARCH FOR DIR INODE
1. read() root_dir inode_disk
2. traverse() dir_entries data block for matching dir name
3. if found dir_entry{name, inode}, populate() parent_dir w inode

IF open(dir)
struct inode *inode = dir_get_inode(dir)
1. return inode

IF open(file)
dir_lookup(dir, file_name, &inode)
1. traverse() dir_entries data block for matching file name
2. if found dir_entry{name, inode}, return inode

int bytes_written = sys_write(fd, "data to write into file", size)

file_write(file_d->file, buffer, size)
1. int fd -> struct file_desc -> struct file -> inode_disk
2. r() w() bitmap for free data block
3. traverse() inode_disk + write() 0s to file_chunks data block (depends on size)
4. traverse() inode_disk + write() "data" to file_chunks data block (depends on size)

BEFORE:
inode bitmap  11000000
inodes        [d a:0 r:3] [f a:1 r:1] [] [] [] [] [] []
data bitmap   10000000
data          [(.,0) (.,0) (y,1)] [0000000000] [] [] [] [] [] []

AFTER:
inode bitmap  11000000
inodes        [d a:0 r:3] [f a:1 r:1] [] [] [] [] [] []
data bitmap   11000000
data          [(.,0) (.,0) (y,1)] [u] [] [] [] [] [] []

```

DEMO open file + write file

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
open(bar)			read		read	read				
				read		read				
					read					
read()					write		read			
					read					
read()					write			read		
					read					
read()					write					read

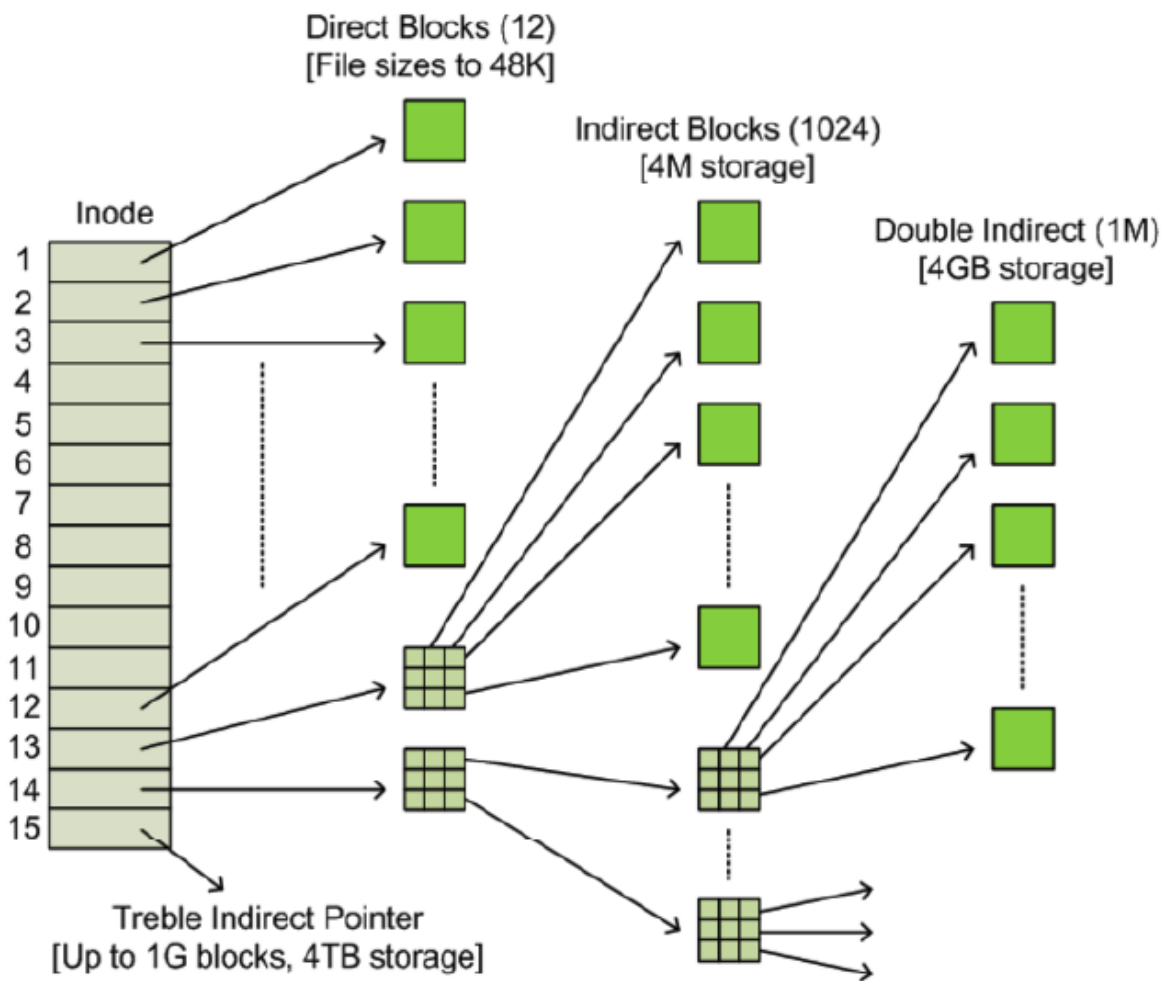
Figure 40.3: File Read Timeline (Time Increasing Downward)

DEMO open() “/foo/bar”

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create (/foo/bar)		read write	read			read				
				read			read			
					read		write			
				write	write					
					read					
write()	read write							write		
					write					
					read					
write()	read write							write		
					write					
					read					
write()	read write									write
					write					

Figure 40.4: File Creation Timeline (Time Increasing Downward)

DEMO create() “/foo/bar”



DEMO inode (USC project 4 slide)

- **USC example (DIFFERENT FROM PINTOS)**

- 1 inode_disk = index() 1 file's data
- inode_disk 1st layer - 12 direct_blocks, 1024 indirect_blocks, 1MB double_indirect_blocks, 1GB treble_indirect_blocks
 - Sector/Block == 4KB (green)
 - Direct == $12 * 4KB == 48KB$
 - Indirect == $1024 * 4KB == 4MB$
 - Double Indirect == $1MB * 4KB == 4GB$

- **Pintos example**

- 1 inode_disk = index() 1 file's data
 - inode_disk == 512 Bytes == $123 * 4 + 4 + 4 + 12$ bytes
 - indirect_inode_disk == 512 bytes == $128 * 4$ bytes
 - file_chunks / dir_entries block == 512 bytes
- inode_disk 1st layer - 123 direct_blocks, 1*128 indirect_blocks, 1*128*128(16384) double_indirect_blocks
 - Sector/Block == 512 Bytes (green)
 - Direct == $123 * 512B == 63KB$
 - Indirect $128 * 512B == 66KB$
 - Double Indirect = $16384 * 512B = 8.3MB$